

Introduction to Neural Networks for Natural Language Processing

Course Overview

Course Overview

- Objectives
- Contents of this Course
- Style & Organization

Objectives

- Introduce the necessary ideas to understand neural networks.
- Allow you to design your own neural architectures.
- Apply neural networks to natural language processing problems.

Contents of this Course

- Four different lessons, each covering a different topic:
 1. Preliminary notions to understand neural networks.
 2. Introduction to neural networks.
 3. Using neural networks for natural language.
 4. Advanced neural networks for structured data.
- Videos, slides and programming notebooks to explore each lesson in detail.

Style & Organization

- The course is meant to be **practical**: what matters is that you learn to use neural networks to solve problems.
- As an introductory course, we focus on an approachable problem: **text classification**.
- To improve the course as other people learn with you, the notebooks are all **open source** and **free on Github**.
- Concepts are explained as **simply as possible**, assuming no previous knowledge other than programming.

Introduction to Neural Networks for Natural Language Processing